

TerraPulse — Demo Run Sheet

Content is user-generated and unverified.

5

DEMO RUN SHEET · PREPARED FOR TOMORROW

TerraPulse, presented to a gold explorer

A working sequence for a 40-minute session, the failure modes worth knowing before you open the laptop, and a question set built to get past politeness.

CommodityGold – full ranking supported

RoomExploration geologist + QP

VehicleObuasi Trend (Demo)

Runtime ~40 min + questions

Before you sleep tonight

Six checks. The first three are the ones that turn a good demo into a broken one, and all three fail silently — nothing warns you until you are already in front of them.

☐ Confirm an AI provider is enabled and holds a key

Admin → AI. If no provider is configured, enabled and keyed, every narration path silently falls back to the offline client: report sections come out as terse deterministic prose and the Copilot answers from canned templates. Nothing errors. It just reads flat, which is exactly the impression you cannot afford with a technical room.

☐ Check the demo tenant's analysis quota and credits

Running an analysis is gated on both an analysis_runs quota and an analysis_run credit. Exhaust either and the Run Analysis button returns an entitlement error mid-session. Seeding the demo project is deliberately not metered, so a healthy demo project tells you nothing about whether a live run will go through.

☐ Re-arm the demo reveal on the browser you'll present from

The first Run Analysis on the demo project is a one-time moment, remembered per device. If you have already used it while testing, clear terrapulse.demo.revealed.<projectId> from that browser's local storage. Otherwise the reveal modal never appears and the button starts a real re-run instead — a slower, less legible version of the same thing.

☐ Seed the demo project now, not live

Seeding krige-interpolates a 1,984-station magnetic survey, normalises 400 geochem samples, runs full scoring and drafts a report. It is the real pipeline, not a fixture — which is the point, but it is not a thing to watch happen in a meeting.

☐ Stage a second, empty project with an AOI already drawn

You want somewhere to ingest one of their files, or to run a genuine analysis if challenged. Imagery, terrain, magnetics and analysis all require a boundary polygon, so draw it in advance — fumbling a polygon live is a bad look for a spatial product.

☐ Set the room up: theme, zoom, tabs

Pick light or dark and stay there. Close Settings → Profile, Settings → Branding and the social login buttons — those hold visible Coming soon placeholders. Browser zoom around 110% so a QP at the far end of the table can read a score.

The run

Ordered so that credibility is earned before capability is claimed. With a QP in the room, provenance has to land before the AI narrative does — show them the machinery first, and the generated prose reads as a summary of something real rather than as the product itself.

00:00

3 min

Frame the problem, laptop closed

State the gap you are addressing before showing anything: evidence arrives in incompatible formats, integration is manual, and the reasoning behind a target lives in someone's head. Ask them to correct you. Their correction is the most useful thing you will hear all day.

00:03

5 min

Open the worked example

Obuasi Trend (Demo) → Map. A prompt offers to reveal the ranked targets; take it. Progress runs, targets appear scored and ranked.

SAY IT PLAINLY

"This project was analysed when it was created — I'm revealing the result, not computing it now. I'll run a real one on live data shortly."

That sentence costs you nothing and buys you the room. A QP's entire professional function is verification; if they later work out that a progress bar was theatre, every number you showed becomes suspect.

go

🔍

⏮

⏪

go

🔍

⏮

⏪

go

🔍

⏮

⏪

go

🔍

⏮

⏪

00:08
7 min

Evidence attribution, target by target

- Targets list — ranked, scored, each above the score floor.
- Open the top target. Score, depth estimate, cell count.
- The evidence drivers: what pushed this target up, and by how much.

This is the part a geologist cannot get from a heat map. Ask directly: *would this attribution change how you'd argue for a hole?*

00:15
5 min

Provenance — the screen for the QP

On the map layer panel, open a layer's "How was this calculated?" drawer: source dataset, processing lineage, kriging parameters, the variogram model actually fitted. Then the geochem thresholds — the 85th, 95th and 99th percentile classes computed from their own population, not fixed cut-offs.

THE QUESTION TO ASK

"What else would you need to see here before you'd sign your name under a target this produced?"

00:20
6 min

Reports, including one generated live

- The demo already carries a drafted Technical Exploration report — open it, show sections and citations.
- Then generate a **Data QA report** live. It runs the qualified-person persona, which is written for the person sitting in front of you.
- Export to PDF. Mention CSV, KML and shapefile for their GIS.

Say that sections are drafted, versioned and editable, and that a new generation creates a new version rather than overwriting — history is preserved. That framing matters to anyone with disclosure obligations.

00:26
6 min

Copilot, grounded in the project

- "What reports have been generated for this project?" — then ask it to summarise one. It reads the actual report text.
- "Show me the magnetics layer" — it changes the map in front of them.
- "Compare the target scores" — it renders a figure from live data.

Point out the citations and that every figure is resolved server-side from the database, so a chart cannot carry an invented number. For this audience that constraint is a feature, not a limitation — say so.

00:32
4 min

The field workflow

Write a qualified-person note on a target. Explain that it holds when the connection drops and syncs on reconnect. Then ask whether their people actually work off-grid — if yes, this is a differentiator worth far more airtime; if no, drop it and take the time back.

00:36
4 min

Their data, on the screen

Offer to ingest a file of theirs into the staged project. Geochem CSV needs a latitude column, a longitude column and at least one numeric element column; geophysics point files need three columns in x, y, value order. If they brought nothing, ask what they would have brought — the answer tells you what to build next.

Known edges

Each of these is a real boundary in the current build. Named first by you, they read as engineering judgement; discovered by them, they read as overselling.

CRITICAL Only gold is ranked

The prospectivity model ranks orogenic gold and nothing else. A copper or lithium project still gets maps, layers, geochemistry and imagery — but no ranked targets. Today's room is gold, so this will not bite during the demo. It will bite the moment they mention the rest of their portfolio.

The honest answer: the weights, lithology favourability and deposit model are tuned for orogenic gold, so ranking is withheld rather than returning numbers from the wrong model. Per-commodity profiles are specified and not yet built. Offer a date, not a maybe.

CRITICAL The demo's Run Analysis reveals precomputed targets

On the demo project only, the progress bar is a client-side reveal of targets computed during seeding. It exists so the "I ran it and got targets" moment takes seconds instead of a minute.

Handling: narrate it as a reveal, then run a genuine analysis on the staged project so they have seen the real thing at real speed.

HIGH No boundary, no anything

Sentinel-2 imagery, terrain, magnetics and prospectivity analysis all require an area-of-interest polygon on the project. Without one they refuse rather than guess.

HIGH Analysis consumes quota and credits

A live run can be refused on entitlement grounds. Confirm headroom before you start and again if you plan more than one run.

HIGH An invited viewer cannot run an analysis

The viewer bundle carries map and reports, not analysis. If you demonstrate collaboration by inviting someone into the tenant live, do not then hand them the Run Analysis button.

NOT BUILT Placeholders to steer around

Social sign-in, password change, and branding accent colour and font are visible but disabled. Nothing breaks if you open them — they simply say "Coming soon", which invites a question you gain nothing by answering.

Getting feedback worth having

People are polite about demos and specific about their own problems. Every question below is built to move them from the first mode to the second. Ask fewer of them, slowly, and write down the exact words.

FORCE A RANKING

Enthusiasm is free. Ordering costs something, so it tells you the truth.

- 1. Of everything you saw, which one thing would you use in the next thirty days?
- 2. Which part would you not use at all — and is that because it's wrong, or because it's not yours to use?
- 3. If we could only finish one of these before you started, which?

FIND WHAT IT REPLACES

A tool that adds a step gets abandoned. A tool that removes one gets budget.

- 1. Walk me through how you do this today — who, which software, how long?
- 2. What would you stop doing if you had this?
- 3. Where does your current process actually hurt — the part you'd complain about unprompted?

SURFACE THE OBJECTION THEY'RE BEING TOO POLITE TO RAISE

Ask for the objection directly or you will hear it months later as silence.

- 1. What would stop your team adopting this, assuming it worked exactly as shown?
- 2. What would your QP need before signing off on a target this produced?
- 3. Who else has to say yes, and what will their objection be?
- 4. Where do you not trust it? Be blunt — that's the useful answer.

TEST THE DATA REALITY

Everything downstream depends on what they actually hold, in what state.

- 1. What survey data do you hold, in what formats, and how clean is it really?
- 2. Is it in a database, on a share drive, or on someone's laptop?
- 3. Historic data — digitised, or still paper and scans?
- 4. What would you never be allowed to upload to a hosted system?

TEST INTENT

The single best predictor of whether this goes anywhere.

- 1. If we gave you access Monday, which project would you load first — and why that one?
- 2. What would have to be true for this to be worth paying for?
- 3. Would you be willing to run a real licence through it with us over two weeks?

Capture it while it's said

Bring someone whose only job is writing. Verbatim quotes, not summaries — a paraphrase loses the specific noun that told you what to build. Fill this in during, not after.

FEATURE	REACTION, IN THEIR WORDS	WOULD USE
Ranked drill targets		1 2 3 4 5
Evidence attribution per target		1 2 3 4 5
Provenance / lineage drawer		1 2 3 4 5
Generated reports + export		1 2 3 4 5
Copilot Q&A over the project		1 2 3 4 5

Offline QP field notes	1	2	3	4	5
Public data fetch (imagery, terrain, magnetics)	1	2	3	4	5
Data ingest from their own files	1	2	3	4	5

Also log, separately: every moment they leaned in, every moment they went quiet, and any question you could not answer. The unanswered questions are your roadmap.

Within 24 hours

1. Send one page: what you heard them say they needed, in their words, and what you will do about each.
2. Ship or schedule the smallest thing they asked for, and tell them the date. Responsiveness is the only credibility a young product has.
3. Name the boundary you hit — if the portfolio question came up, put the multi-commodity position in writing before they ask twice.
4. Propose the concrete next step you already floated: one real licence, two weeks, a named contact on each side.

Run sheet reflects the build on branch feat/copilot-chat-ux-report-map · verify the pre-flight items against the environment you actually present from

